

OpenAMP – Intro

Feb 2016

Tomas Evensen, Xilinx

Intro

With today's sophisticated SoCs there is often a need to integrate multiple runtime environments with multiple operating systems. This raises a lot of issues, such as:

- Lifecycle issues
 - Boot order, start one OS from another, tear down OS, reboot, ...
- Communication
 - Message passing, ...
- Resource handling
 - Memory, devices, interrupts, ...
- Power management
- ...

OpenAMP aims to address these and other issues in a standardized way, both through an open source project and through standardization by MCA.

What is Needed to Be Able to Mix OSes?

- A standard “protocol” so OSes can interact without dependencies
 - On a given HW, any OS can interact with any other OS
 - Without special adaptation
 - On a shared memory system this is a set of data structures and conventions
 - I.e. the ring buffers in *virtio*
- A standard set of APIs for OS interactions
 - A low level API that abstracts underlying OS and HW
 - A set of lifecycle APIs
 - Messaging APIs
 - Other potential features
 - Proxy capabilities to make remote OS look like Linux process
 - Remote procedure calls, power management, device configuration, debug, ...
- Upstream Linux support for protocol and APIs
 - Linux is increasingly becoming the main OS in a multi-OS system
- Open Source implementation
 - Quickest way to adoption
 - Standardization by reference implementation

What is OpenAMP?

- OpenAMP standardizes how Operating Systems interact
 - In particular between Linux and RTOS/bare-metal
 - In particular in a multicore heterogeneous systems
 - Includes:
 - Shared memory protocol for OS interactions (virtio)
 - Lifecycle APIs to start/stop/... other OSES (rproc)
 - Communication APIs to share data (rpmmsg)
 - More to come
- Both a standardization effort and an open source project
- Guiding principles
 - Open Source implementations for Linux and RTOSes
 - Prototype and prove in open source before standardizing
 - Business friendly APIs and implementations to allow proprietary solutions

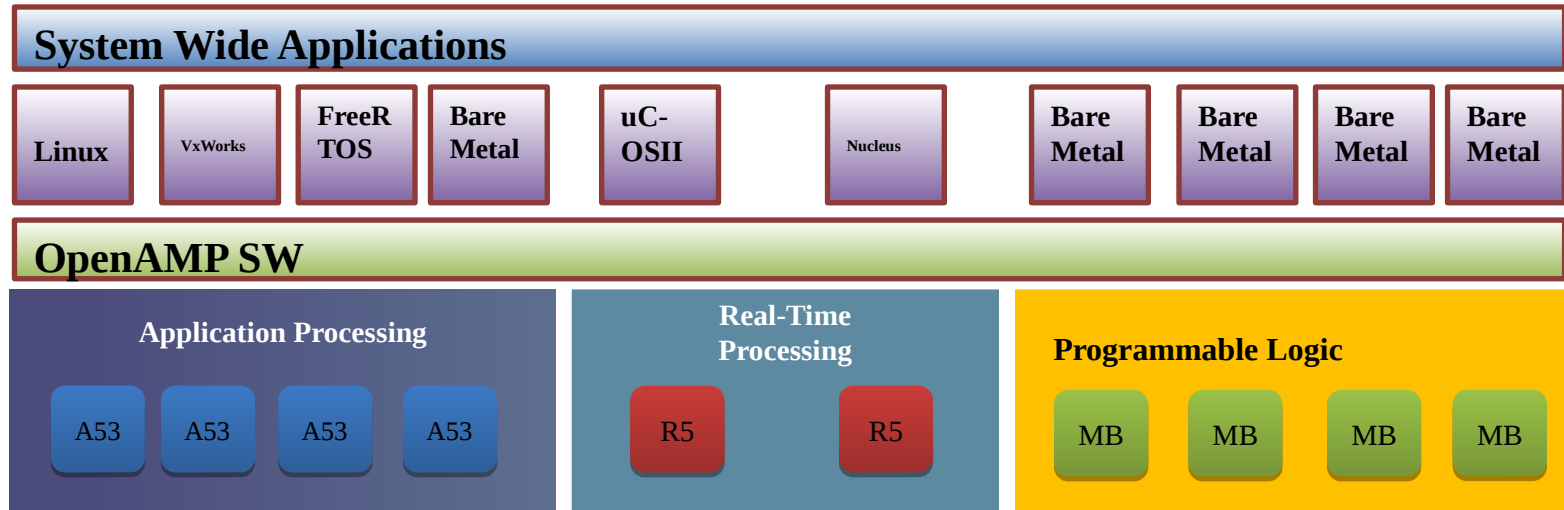
Open Source and Standardization

- Same players but different focus
 - Value in separating standardization and implementation
- Open Source project
 - Focus on implementation and testing new ideas
 - Today driven by Mentor Graphics and Xilinx
 - Also used by Freescale, Wind River and Micrium
 - Based on existing open source technologies in Linux
 - Shared memory protocol – virtio, Lifecycle APIs – remoteproc, Messaging – rpmsg
 - Clean-room implementation for non-Linux kernel code
 - Multiple licenses (GPL for Linux kernel, BSD for rest)
 - Proxy technologies to emulate Linux processes
 - More in future, please join!
- MCA OpenAMP Working Group
 - Focus on formally standardizing APIs used by open source project
 - Allows for proprietary solutions like hypervisors, certified systems, commercial OSes, ...
 - Discuss future directions
 - Implement in open source – standardize when proven

Working Group Mission and Structure

- Mission statement (strawman):
 - “To formally standardize the interactions between Operating Systems in an Asymmetric Multiprocessing (AMP) multicore environment by adopting APIs used in Open Source solutions.”
- Structure:
 - Chair: Tomas Evensen – Xilinx
 - Vice Chair: Felix Baum – Mentor
 - Secretary: Tina Hutchens – Xilinx
 - Participants (Feb 2016) – Xilinx, Mentor, Feescale/NXP, TI, Wind River, Micrium, Express Logic, Airbus, Siemens
- Meetings:
 - Weekly meetings Thu 9 AM Pacific Time

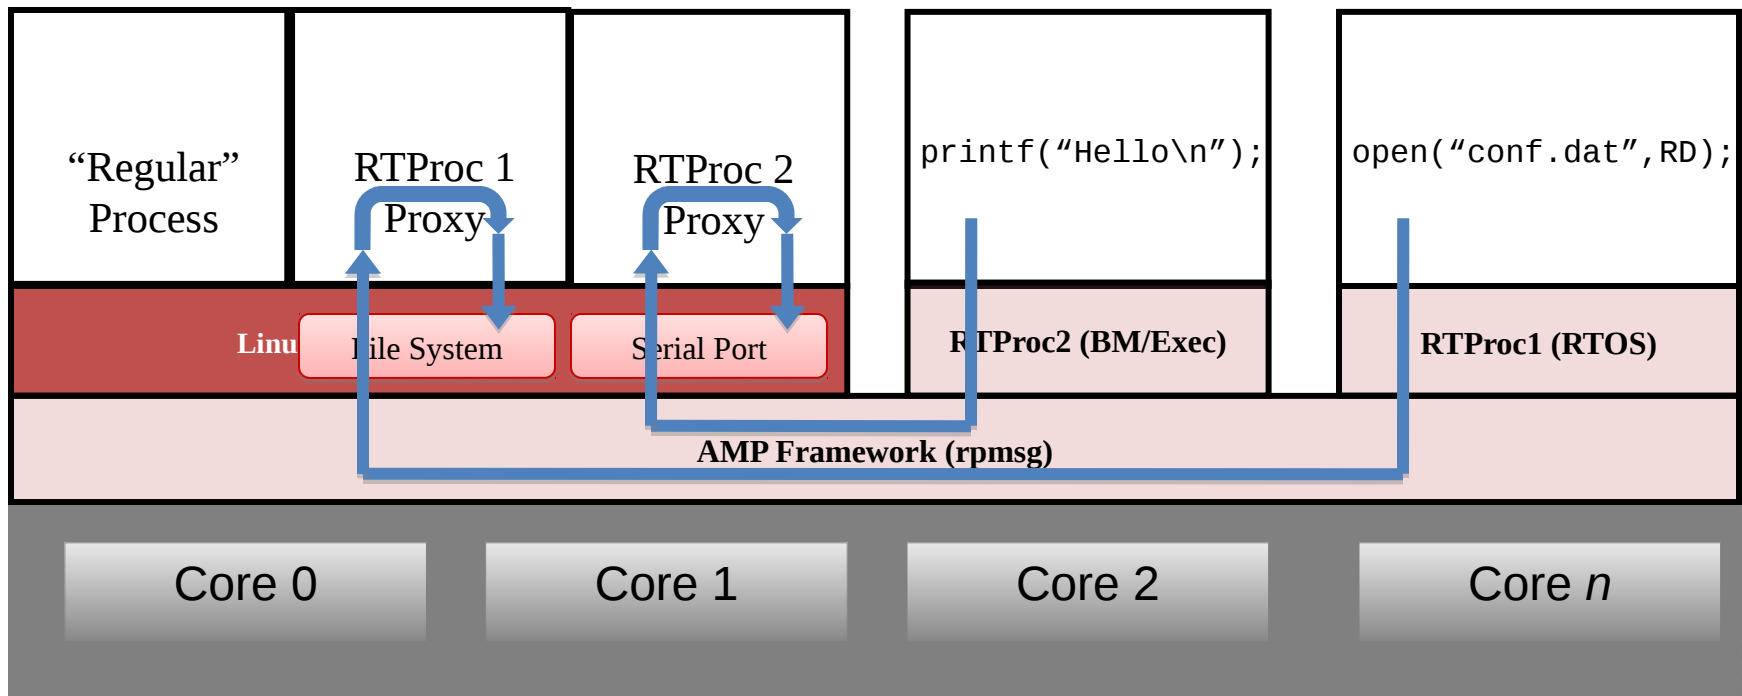
OpenAMP SW Simplifies Heterogeneous Systems



- Provides a Layer for Applications
 - Standard API's that allow applications to be ported across processors and operating systems, including hypervisors
- System Development
 - Provides a wide range of capabilities needed to deploy applications across asymmetric computing elements
- Inter-OS & Inter Processor Communication
 - Send messages back and forth
- OS Management
 - Provides booting/rebooting of processors
- Linux proxy processes
- More features in the future

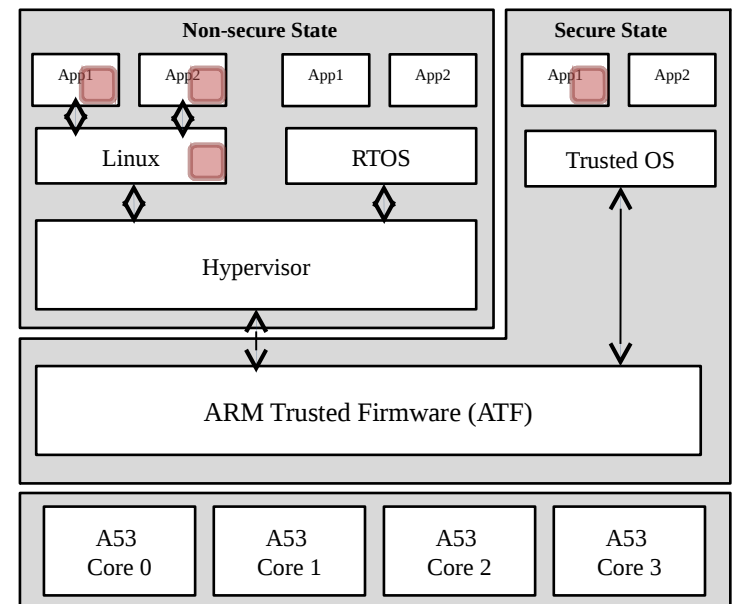
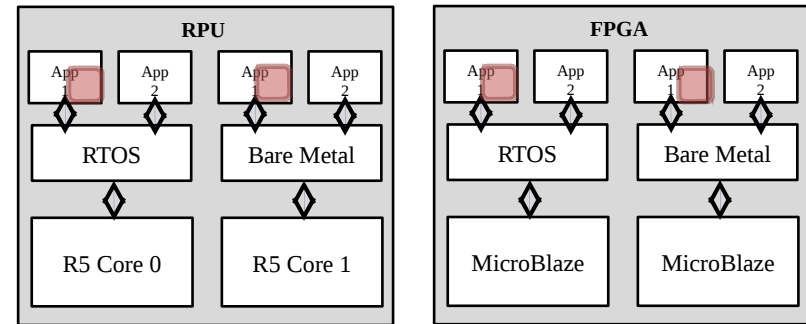
RT Process – Make AMP Look Like Linux

- Use standard Linux process concept to manage other OS
 - Compile application meant to run on RTOS linking to RTOS and OpenAMP libraries
 - Put application executable (ELF file) in the Linux file system
 - Boot Linux as normal on main processors
 - Start application through special OpenAMP launch program, runs in supervisor mode
 - Handles select system(i.e. open/close/read/write/...) calls through RPC
 - Terminate RT process like any other process




OpenAMP Use Cases for Zynq MPSoC

- OpenAMP applications can run in many places
 - A53, R5, MicroBlaze cores
 - In Linux kernel, Linux user space, RTOS, Bare Metal
 - OS's can run directly on core(s), in hypervisor, in trusted mode, ...
- A lot of different operating systems and hypervisors
 - Linux is used in majority of use cases
 - Most free and commercial RTOS's are being used
 - Bare metal (no OS) is common on smaller cores
- Communication typically through shared memory
 - Though dedicated HW often used for data plane
 - Interprocess “kick” done in various ways
 - Inter Processor Interrupt (IPI), kernel call, hyper call, ...
- Boot order varies
 - “Independent” boot
 - Through uboot, hypervisor, ...
 - OpenAMP master starts slave
 - “Big” starts “small” (e.g. Linux app on A53 starts RTOS on R5)
 - “Small” starts “big” (e.g. Safe RTOS starts Linux)
- Safety and Security issues common
 - Affects boot order, messaging implementation, ...



 - Examples of OpenAMP applications

Evolving OpenAMP Architecture (Current)



Virtio – a communication protocol typically using shared memory and IPI

rpmmsg – a set of messaging APIs:


```
int rpmmsg_send(struct rpmmsg_channel *rpdev, void *data, int len);
int rpmmsg_sendto(struct rpmmsg_channel *rpdev, void *data, int len, u32 dst);
int rpmmsg_send_offchannel(struct rpmmsg_channel *rpdev, u32 src, u32 dst,
int rpmmsg_trysend(struct rpmmsg_channel *rpdev, void *data, int len);
int rpmmsg_trysendto(struct rpmmsg_channel *rpdev, void *data, int len, u32 dst);
...
```

Remoteproc – a set of APIs dealing with startup/shutdown:

```
int rproc_boot(struct rproc *rproc);
void rproc_shutdown(struct rproc *rproc);
struct rproc *rproc_get_by_phandle(phandle phandle);
```

Proxy – a Linux user process that enables a remoteproc instance act similar to a Linux process

Evolving OpenAMP Architecture (Available)



Linux user-space – An efficient implementation in user-space utilizing UIO

Evolving OpenAMP Architecture (Next)

Commercial AMP Solution Space
(Mentor, WindRiver, Micrium, ...)

OpenAMP

virtio

rpmsg

remoteproc

proxy

MCAPI

...

Platform

Zynq

ZynqMP-R5

ZynqMP-A53

MicroBlaze

...

OS Adaption

Atomics

Locks

Shmem

I/O Mem

Bus

Device

IRQ

DMA

...

Linux

Generic (bare metal)

RTOS

UIO

VFIO

LibSysfs

LibHugeTLBFS

...

Proposed Working Group Tasks

- Define “North-bound” interfaces
 - Rpmmsg, remoteproc
 - Discuss future enhancements
 - E.g. attach to running OS, communicate available resources, ...
- Define “South-bound” interfaces
 - OS and HW abstraction interfaces
- Discuss MCAPI implementation on top of South-bound interfaces
 - Start with subset?